

MinCaml:

A Simple and Efficient Compiler for a Minimal Functional Language

Eijiro Sumii
Tohoku University



Highlights

"Simple and efficient compiler
for a minimal functional language"

- Only **2000 lines** of OCaml
- Efficiency comparable to OCamlOpt and GCC for several applications
 - **Ray tracing**, Huffman encoding, etc.
- Used for **undergraduates** in Tokyo since 2001

Outline of This Talk

- Pedagogical background
- Design and implementation of MinCaml
- Efficiency

Computer Science for Undergraduates in Tokyo

- Liberal arts (1.5 yr)
 - English, German/Chinese/French/Spanish, mathematics, logic, physics, chemistry, ...
 - Computer literacy, CS introduction, Java programming, data structures
- CS major (2.5 yr) [~30 students/yr]
 - Algorithms, OS, architecture, ...
 - SPARC assembly, C, C++, Scheme, OCaml, Prolog

Programming Languages for CS Major in Tokyo

- PL labs (63 hr)
 - Mini-Scheme interpreter in Scheme,
 - Mini-ML interpreter in OCaml,
 - Othello/Reversi competition in OCaml, etc.
- Compiler lecture (21 hr)
 - Parsing, intermediate representations, register allocation, garbage collection, ...
- PL theory lectures (42 hr)
 - λ -calculus, semantics, type theory, ...

CPU/Compiler Labs (126 hr)

- CPU lab
 - Design and implement original CPUs by using VHDL and FPGA
 - Compiler lab
 - Develop compilers for the original CPUs
 - ✓ MinCaml is used here!
- ⇒ Compete by the speed of ray tracing
(5-6 students per group)

Extension board

FPGA board

XC2V1000
FPGA


SSRAM


SDRAM


USB port

Configuration
cable

Special thanks to Hiraki laboratory







How is MinCaml Used?

- Students are given high-level descriptions of MinCaml
 - in Japanese and pseudo-code
- Each group is required to implement them
- Every student is required to solve small exercises
 - such as hand compilation

Outcome (1/2)

Students liked ML!

- Implemented polymorphism (like MLton), garbage collection, inter-procedural register allocation, etc. **without being told**
- Started a portal site (www.ocaml.jp) with Japanese translations of the OCaml manual **without being told**

Outcome (2/2)



"Outsiders" are also using MinCaml

- Somebody anonymous wrote a comprehensive commentary on MinCaml
- Ruby hackers organized an independent seminar to study MinCaml
- Prof. Asai is using MinCaml in Ochanomizu University

Outline of This Talk

- Pedagogical background
- Design and implementation of MinCaml
- Efficiency

Goals



- As simple as possible

but

- Able to efficiently execute non-trivial applications (such as ray tracing)

MinCaml: The Language

- Functional: no destructive update of variables (cf. SSA)
- Higher-order
- Call-by-value
- Impure
 - Input/output
 - Destructive update of arrays
- Implicitly typed
- Monomorphic

Syntax (1/2)

M, N (expressions) ::=

c

$op(M_1, \dots, M_n)$

$\text{if } M \text{ then } N_1 \text{ else } N_2$

$\text{let } x = M \text{ in } N$

x

$\text{let rec } x \ y_1 \dots y_n = M_1 \text{ in } M_2$

$M \ N_1 \dots N_n$ (no partial application)

...

(cont.)

Syntax (2/2)

M, N (expressions) ::=

...

(M_1, \dots, M_n)

$\text{let } (x_1, \dots, x_n) = M \text{ in } N$ (cf. $\#_i M$)

$\text{Array.create } M \ N$

$M.(N)$

$M_1.(M_2) \leftarrow M_3$

$()$

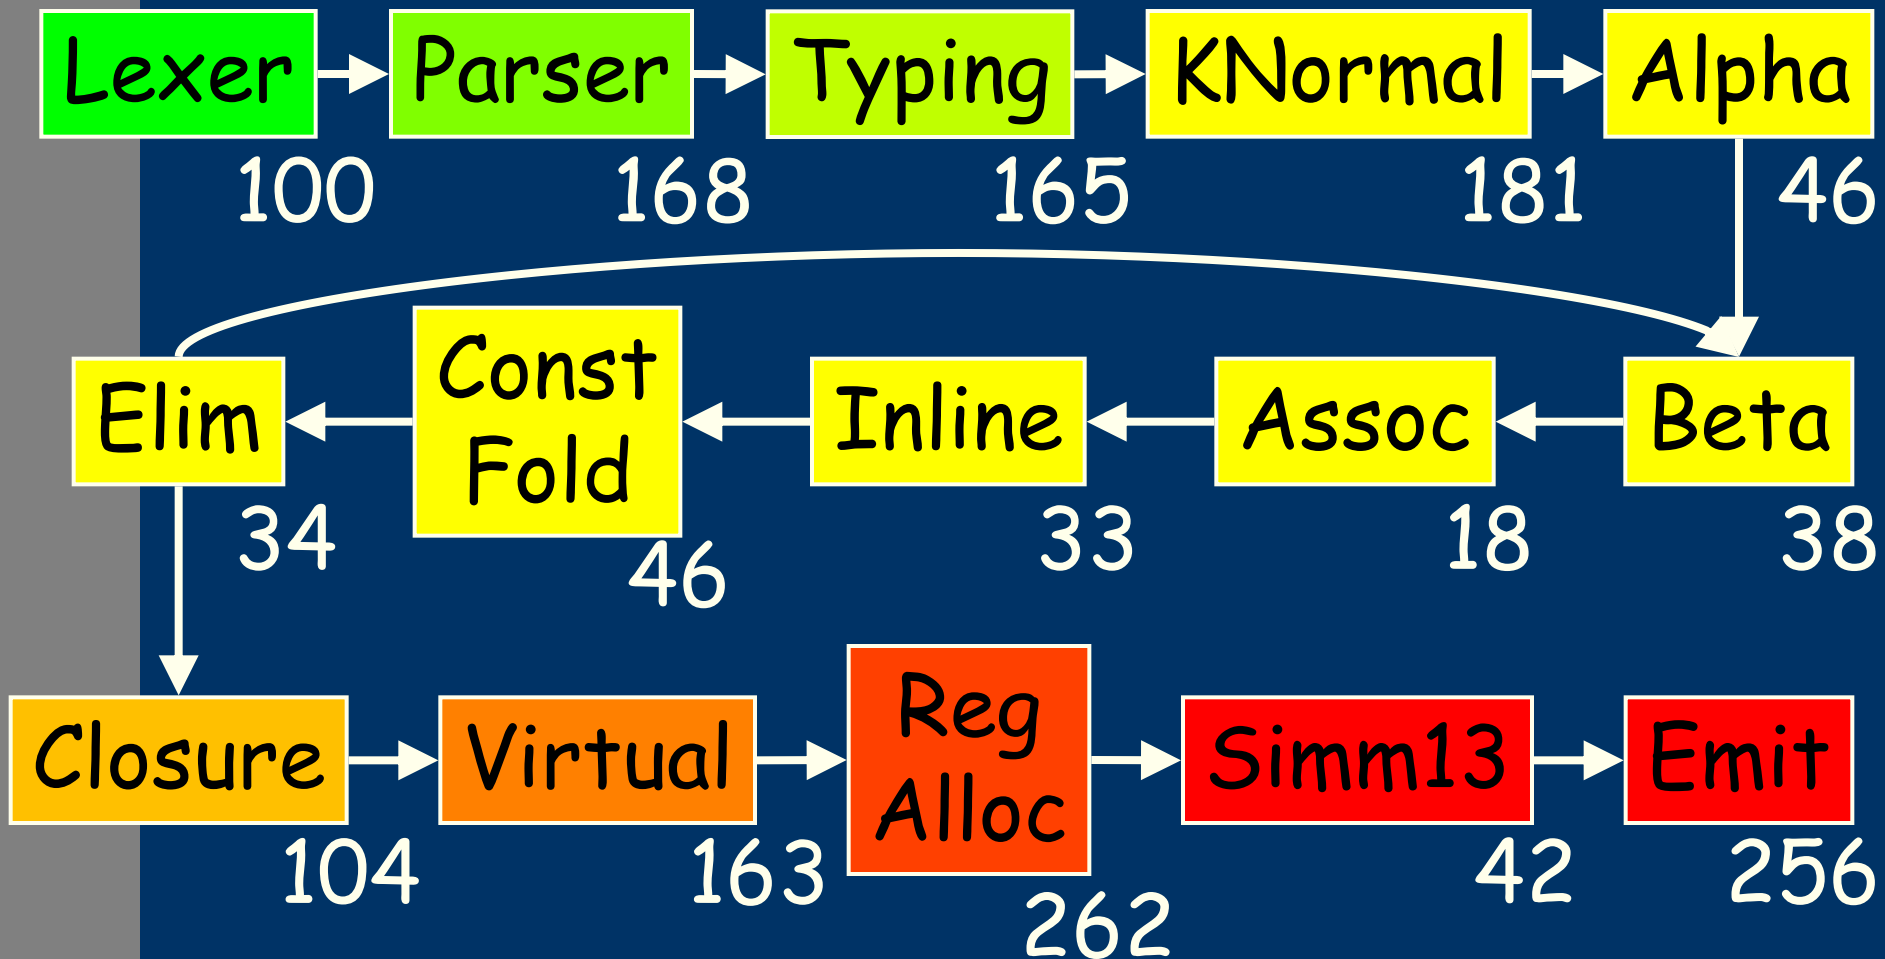
Literally implemented as
ML data type `Syntax.t`

Everything else is omitted!

- Array boundary checking (easy)
- Garbage collection
- Data types and pattern matching
- Polymorphism
- Exceptions
- Objects etc.

Optional homework
(≥ 2 compulsory from this year)

MinCaml: The Compiler



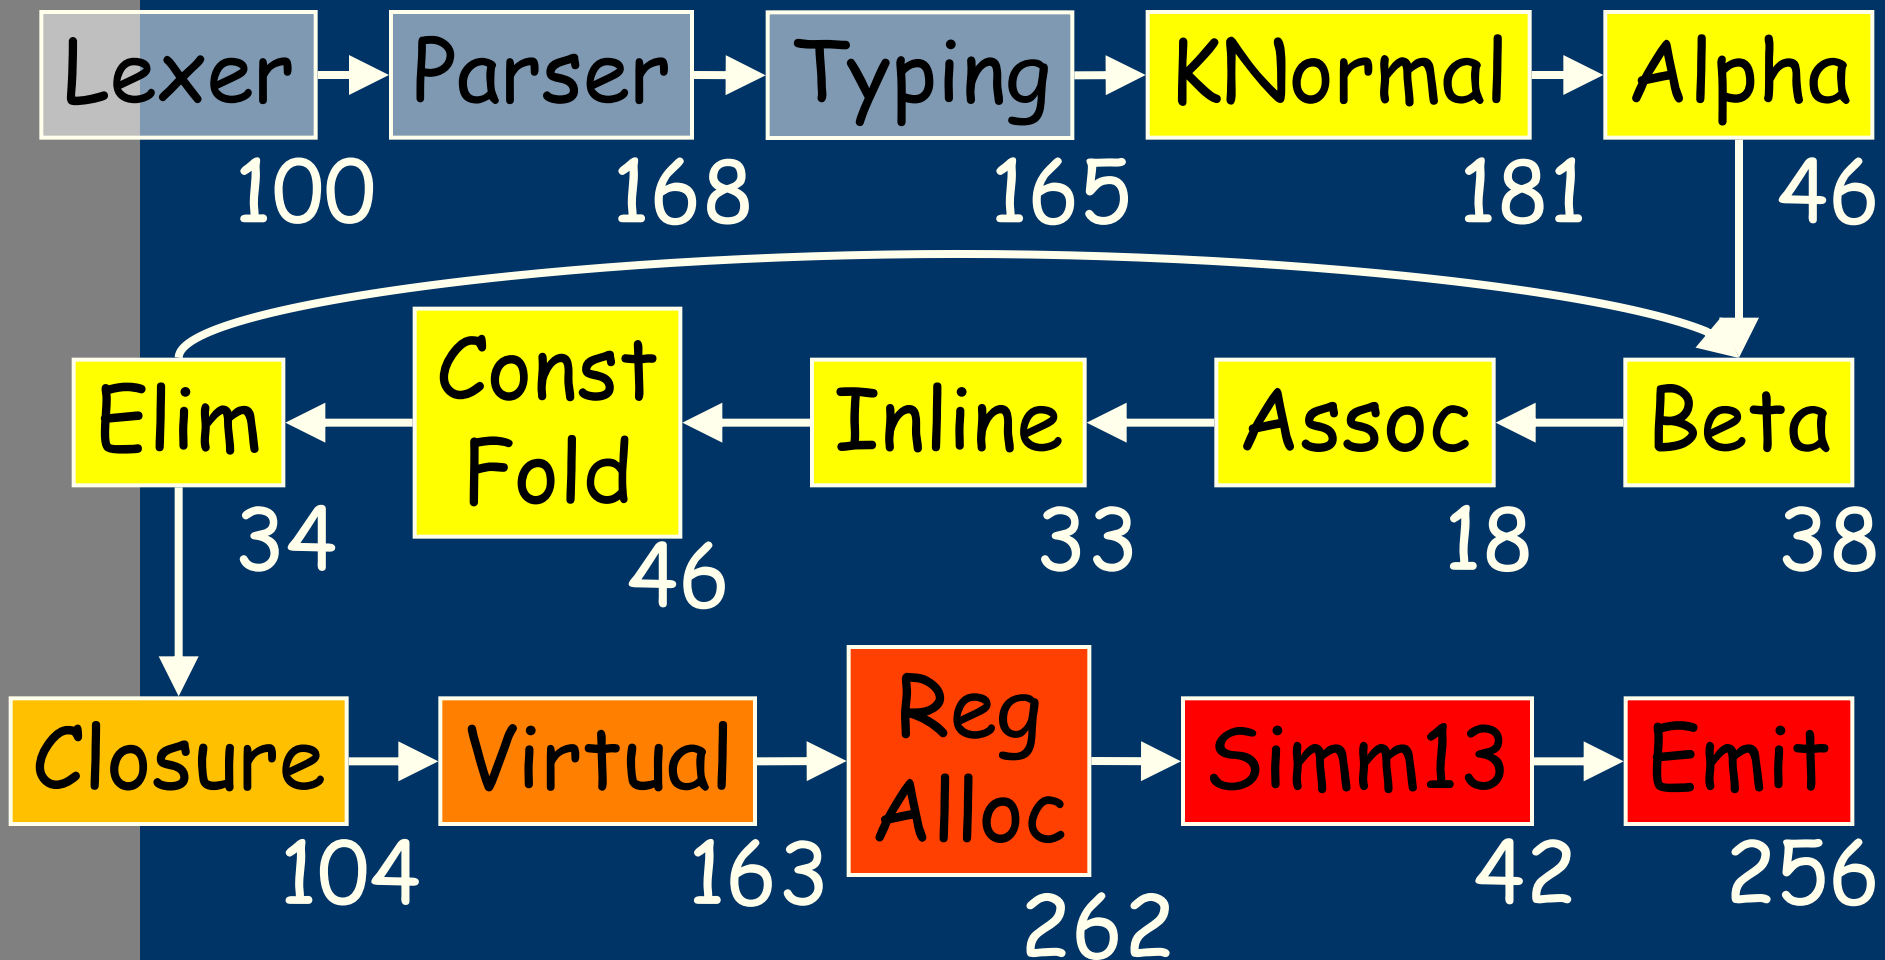
Lexing and Parsing

- Written in OCamlLex and OCamlYacc
 - Given by the instructor
 - Algorithms are out of scope
- Cf. packrat parsing [Ford 2002]

Type Inference

- Based on standard unification using ML references
 - Let-polymorphic version is already taught in PL lab
- Free variables are treated as external functions (or arrays)
 - "Principal typing" [Jim 96] is automatically inferred

MinCaml: The Compiler



K-Normalization

$$a + b + c * d$$


```
let tmp1 = a + b in  
let tmp2 = c * d in  
tmp1 + tmp2
```

- Nesting is allowed

$\text{let } x = (\text{let } y = M_1 \text{ in } M_2) \text{ in } M_3$

- Simplifies the normalization and inlining
- Cf. A-normalization by CPS

Syntax of K-Normal Form

$M, N ::=$

c

$\text{op}(x_1, \dots, x_n)$

$\text{if } x \text{ then } M_1 \text{ else } M_2$

$\text{let } x = M \text{ in } N$

x

$\text{let rec } x \ y_1 \dots y_n = M_1 \text{ in } M_2$

$x \ y_1 \dots y_n$

\dots

Implemented as `KNormal.t`

Algorithm of K-Normalization: Pseudo-Code Given to Students

$K : \text{Syntax.t} \rightarrow \text{KNormal.t}$

$K(c) = c$

$K(\text{op}(M_1, \dots, M_n)) =$

let $x_1 = K(M_1)$ in ... let $x_n = K(M_n)$ in
op(x_1, \dots, x_n)

$K(\text{if op}(M_1, \dots, M_n) \text{ then } N_1 \text{ else } N_2) =$

let $x_1 = K(M_1)$ in ... let $x_n = K(M_n)$ in
if op(x_1, \dots, x_n) then $K(N_1)$ else $K(N_2)$

$K(\text{let } x = M \text{ in } N) = \text{let } x = K(M) \text{ in } K(N)$

$K(x) = x$ etc.

α -Conversion (Another Example of Pseudo-Code)

$\alpha : \text{KNormal.t} \rightarrow \text{Id.t Map.t} \rightarrow \text{KNormal.t}$

$$\alpha(c)_\rho = c$$

$$\alpha(\text{op}(x_1, \dots, x_n))_\rho = \text{op}(\rho(x_1), \dots, \rho(x_n))$$

$$\alpha(\text{if } x \text{ then } N_1 \text{ else } N_2)_\rho =$$

$$\text{if } \rho(x) \text{ then } \alpha(N_1)_\rho \text{ else } \alpha(N_2)_\rho$$

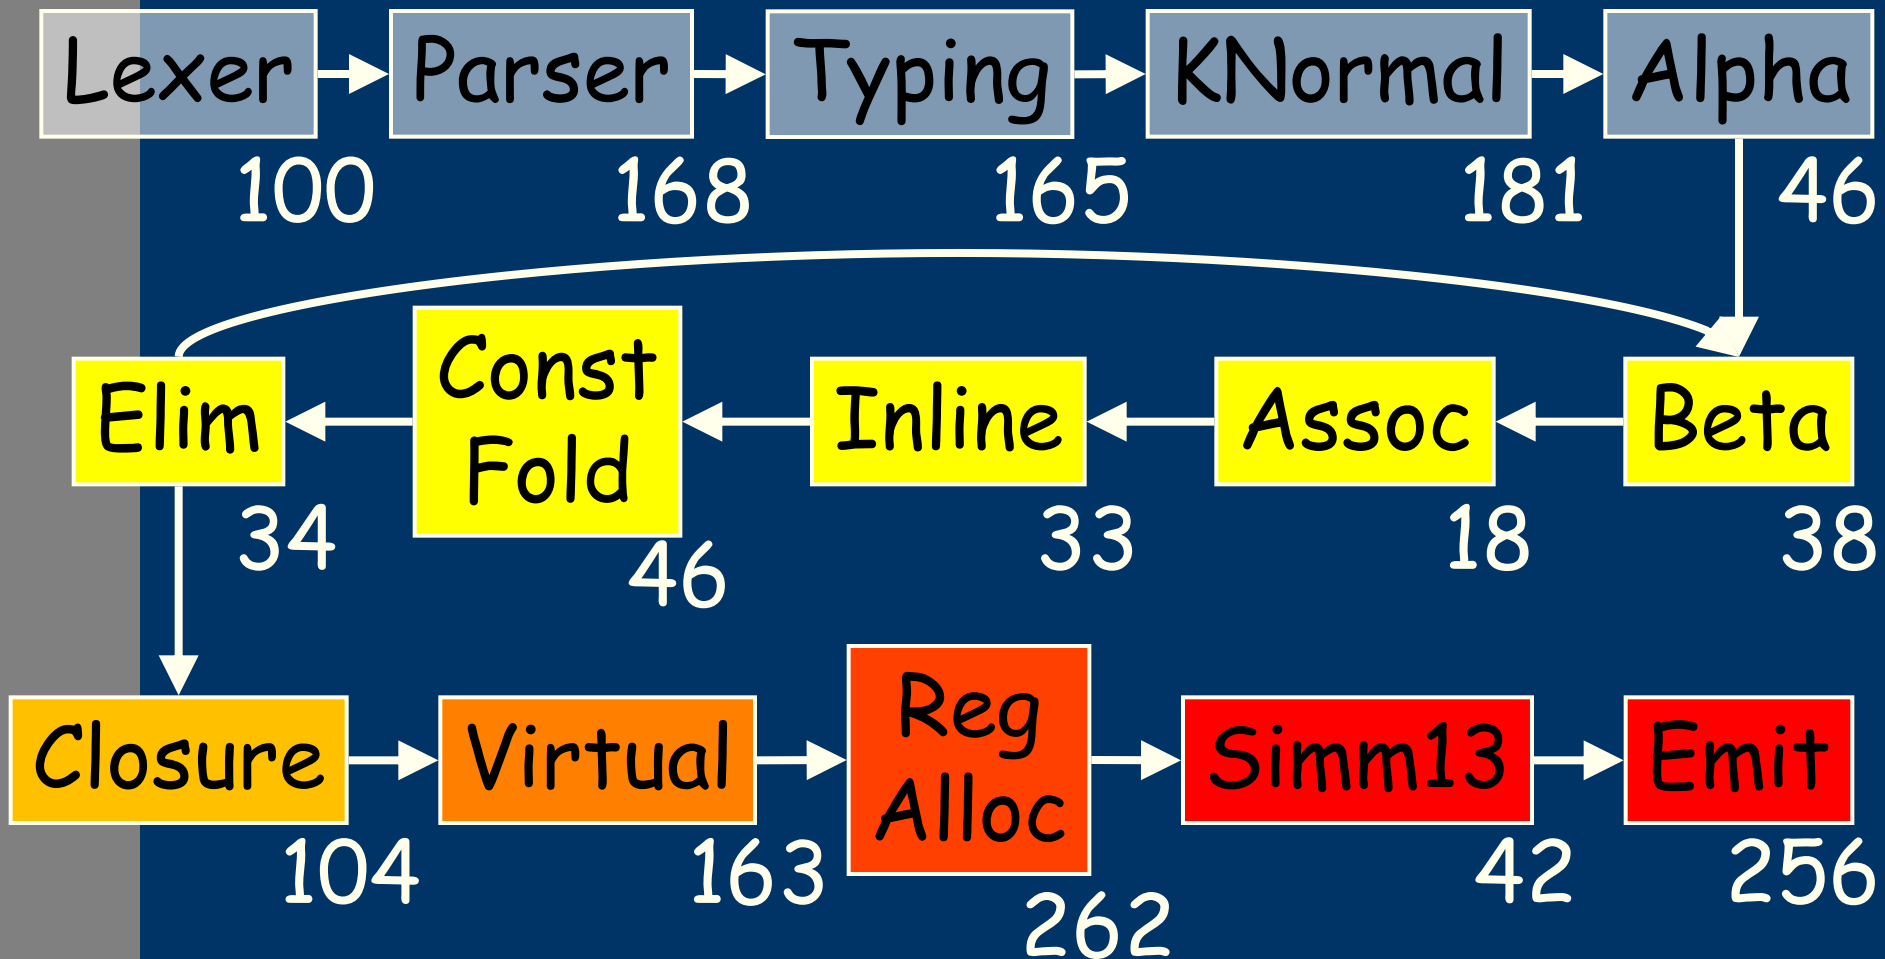
$$\alpha(\text{let } x = M \text{ in } N)_\rho = \quad (x' \text{ fresh})$$

$$\text{let } x' = \alpha(M)_\rho \text{ in } \alpha(N)_\rho[x \rightarrow x']$$

$$\alpha(x)_\rho = \rho(x)$$

etc.

MinCaml: The Compiler



β -Reduction

$$\text{let } x = y \text{ in } M \Rightarrow [y/x]M$$

- Pseudo-code (similar to previous examples) is left as an exercise

Nested "Let" Reduction

$$\text{let } y = (\text{let } x = M_1 \text{ in } M_2) \text{ in } M_3$$
$$\Downarrow$$
$$\text{let } x = M_1 \text{ in let } y = M_2 \text{ in } M_3$$

- Resembles A-normalization, but does not expand "if"

$$C[\text{if } M \text{ then } N_1 \text{ else } N_2]$$
$$\Rightarrow \text{if } x \text{ then } C[N_1] \text{ else } C[N_2]$$

Inlining

Inlines all "small" functions

- Includes recursive ones
- "Small" = less than a constant size
 - User-specified by "-inline" option
- Repeat for a constant number of times
 - User-specified by "-iter" option

Constant Folding and Unused Variable Elimination

let x = 3 in let y = 7 in x + y



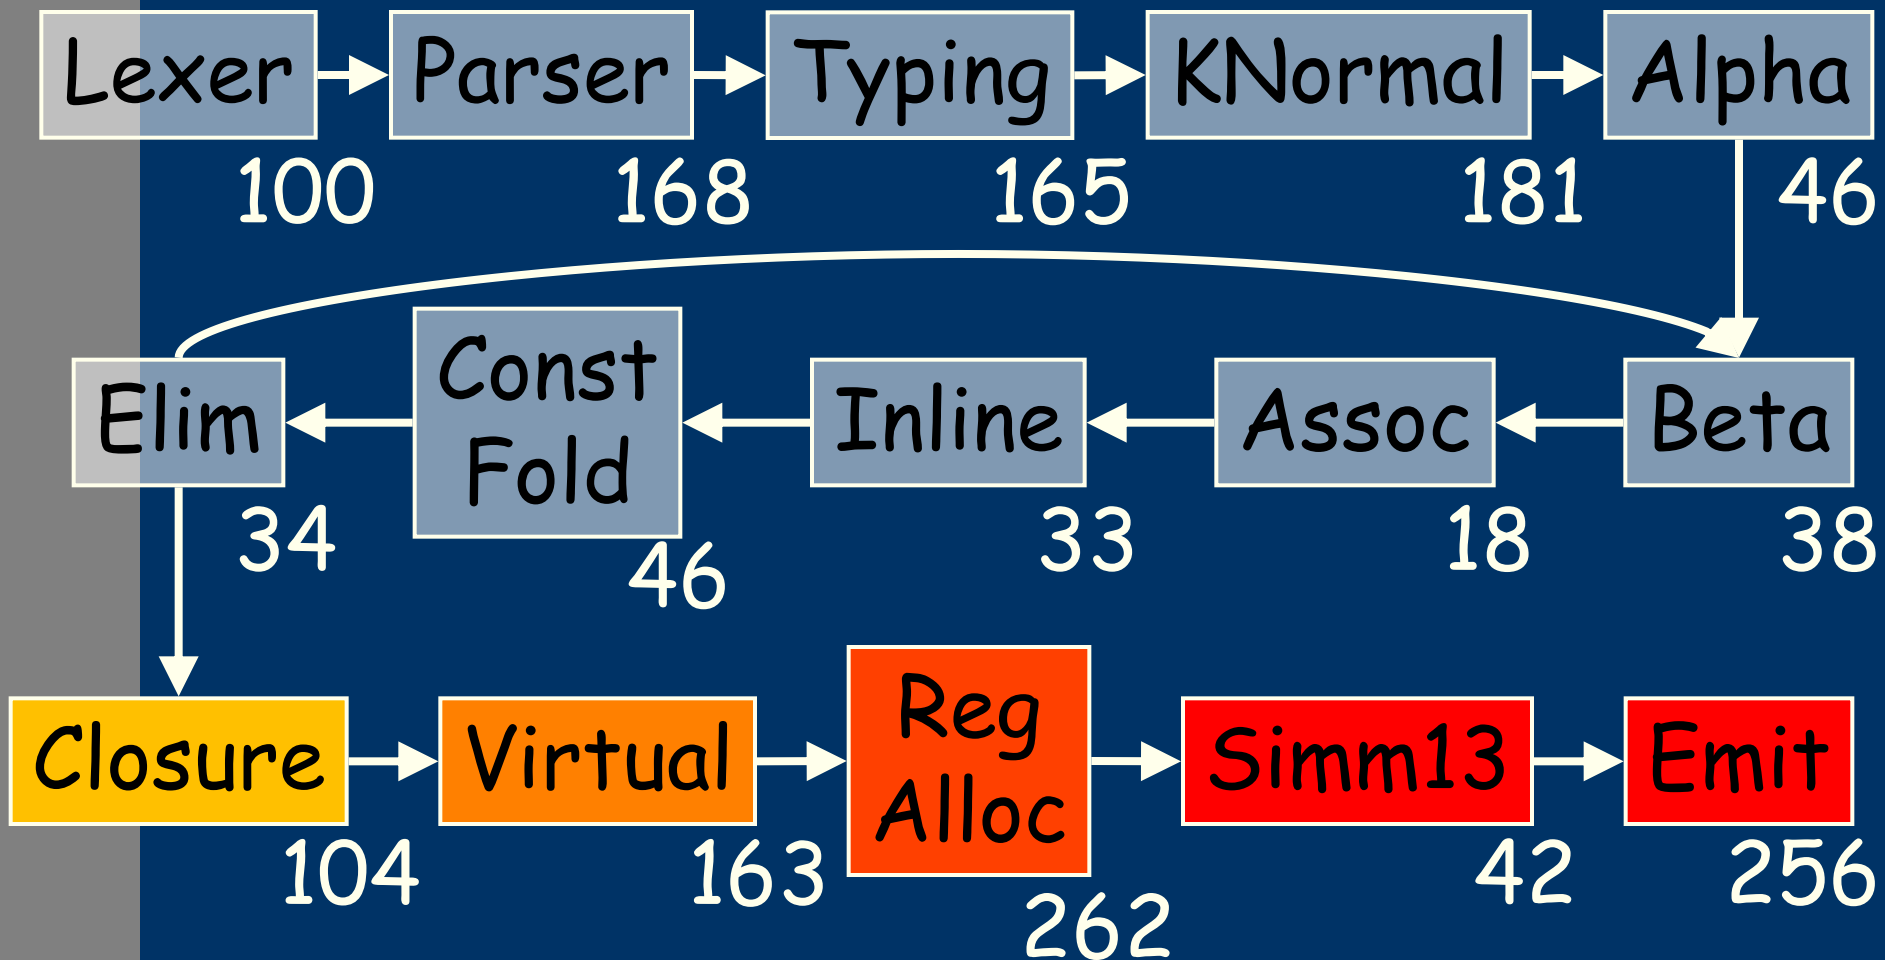
let x = 3 in let y = 7 in 10



10

Effective after inlining

MinCaml: The Compiler



Closure Conversion

Local function definitions (**let rec**)
+ function applications



Top-level function definitions
+

- Closure creations (**make_closure**)
- Closure applications (**apply_closure**)
- Known function calls (**apply_direct**)

Example 1:

Closure Creation/Application

```
let x = 3 in  
let rec f y = x + y in  
f 7
```



```
let rec ftop [x] y = x + y ::
```

```
let x = 3 in  
make_closure f = (ftop, [x]) in  
apply_closure f 7
```

Example 2: Known Function Call

```
let rec f x = x + 3 in  
(f, f 7)
```



```
let rec ftop [] x = x + 3 ;;
```

```
make_closure f = (ftop, []) in  
(f, apply_direct f 7)
```

Example 3:

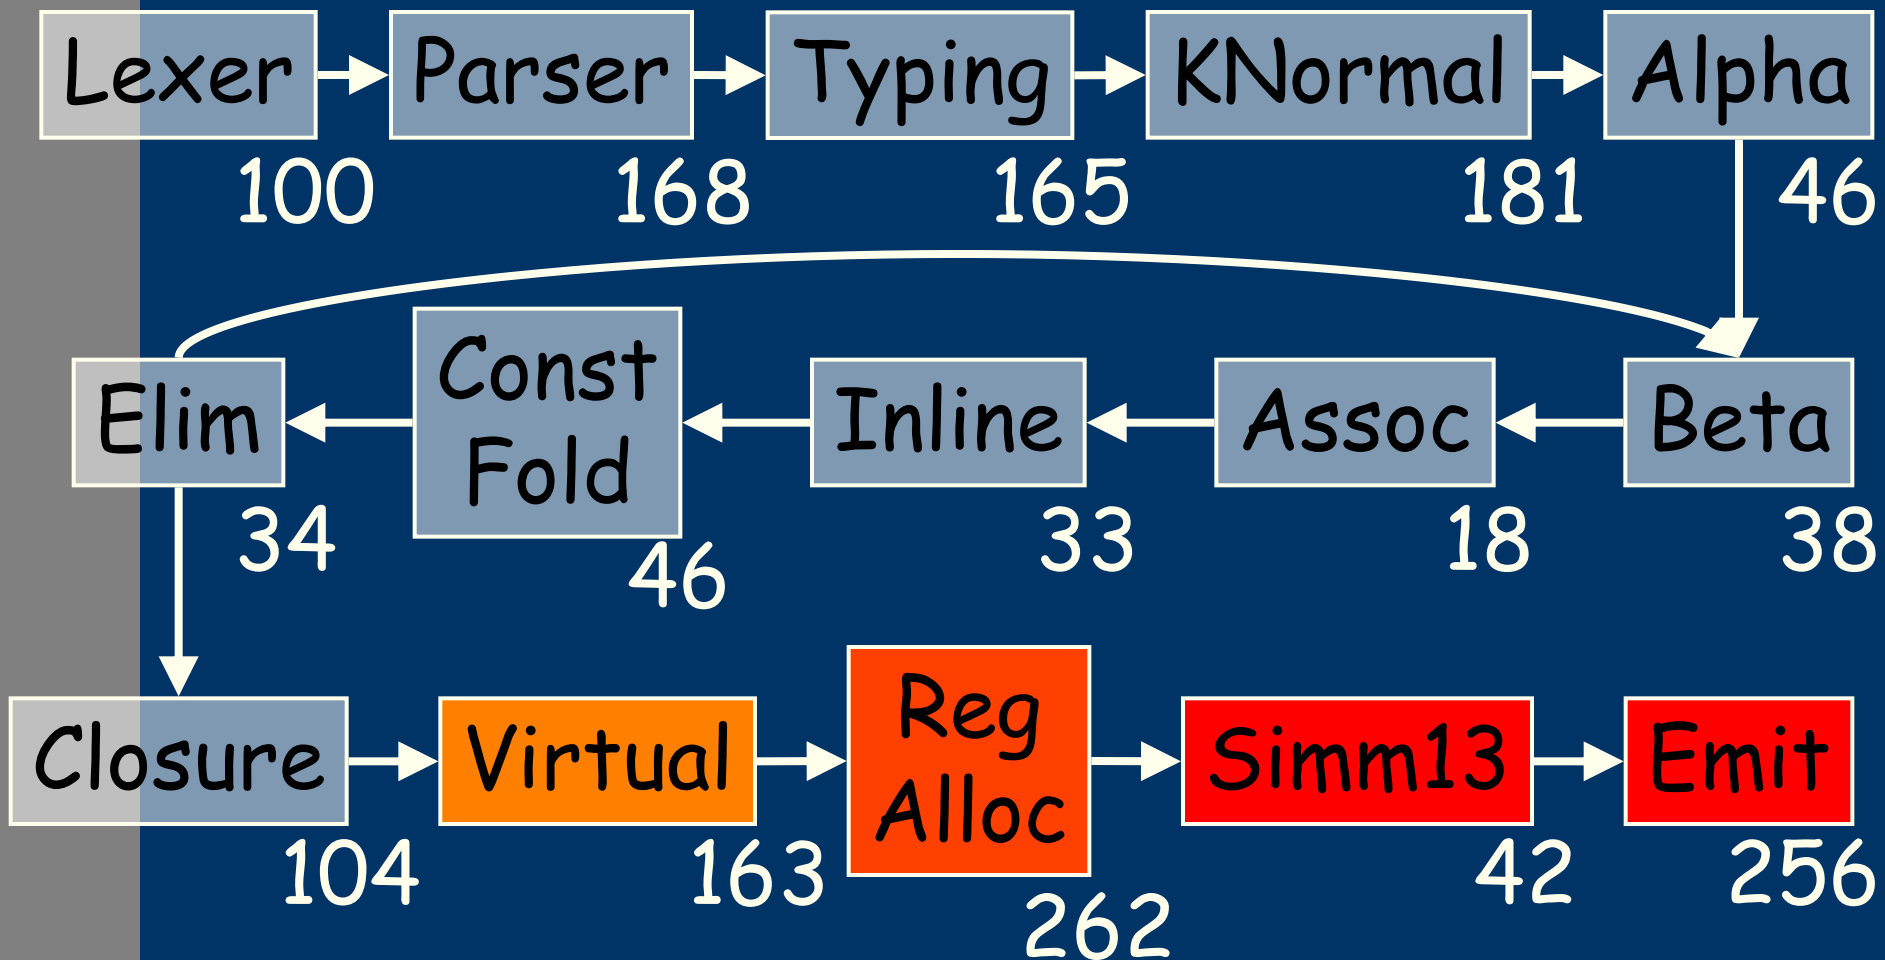
Unused Closure Elimination

```
let rec f x = x + 3 in  
f 7
```



```
let rec ftop [] x = x + 3 ;;  
apply_direct f 7
```


MinCaml: The Compiler



Virtual Machine Code Generation

SPARC assembly with:

- Infinite number of registers/variables
- Top-level function definitions and calls (`call_closure`, `call_direct`)
- Conditional expressions (`if`)

Tuple creations/accesses
and closure creations are
expanded to stores and loads

Register Allocation

Greedy algorithm with:

- Look-ahead for targeting

let $x = 3$ in let $y = 7$ in $f\ y\ x$

\Rightarrow let $r_2 = 3$ in let $r_1 = 7$ in $f\ r_1\ r_2$

- Backtracking for "early save"

let $x = 3$ in

...; $f\ ()$; ...; $x + 7$

\Rightarrow let $r_1 = 3$ in

save(r_1 , x); ...; $f\ ()$; ...; restore(x , r_2); $r_2 + 7$

13-Bit Immediate Optimization

- Specific to SPARC
- "Inlining" or "constant folding" for integers from -4096 to 4095

```
set 123, %r1  
add %r1, %r2, %r3
```



```
add %r2, 123, %r3
```

Assembly Generation

Lengthy (256 lines)
but easy

- Tail call optimization
- Stack map computation
- Register shuffling
 - Somewhat tricky but short (11 lines)

Outline of This Talk



- Pedagogical background
- Design and implementation of MinCaml
- Efficiency


Environment

- Machine: Sun Fire V880
 - 4 Ultra SPARC III 1.2GHz
 - 8 GB main memory
 - Solaris 9
- Compilers:
 - MinCaml (32 bit, -iter 1000 -inline 100)
 - OCamlOpt 3.08.3 (32 bit, -unsafe -inline 100)
 - GCC 4.0.0 20050319 (32 bit and 64 bit, -O3)
 - GCC 3.4.3 (32 bit "flat model", -O3)

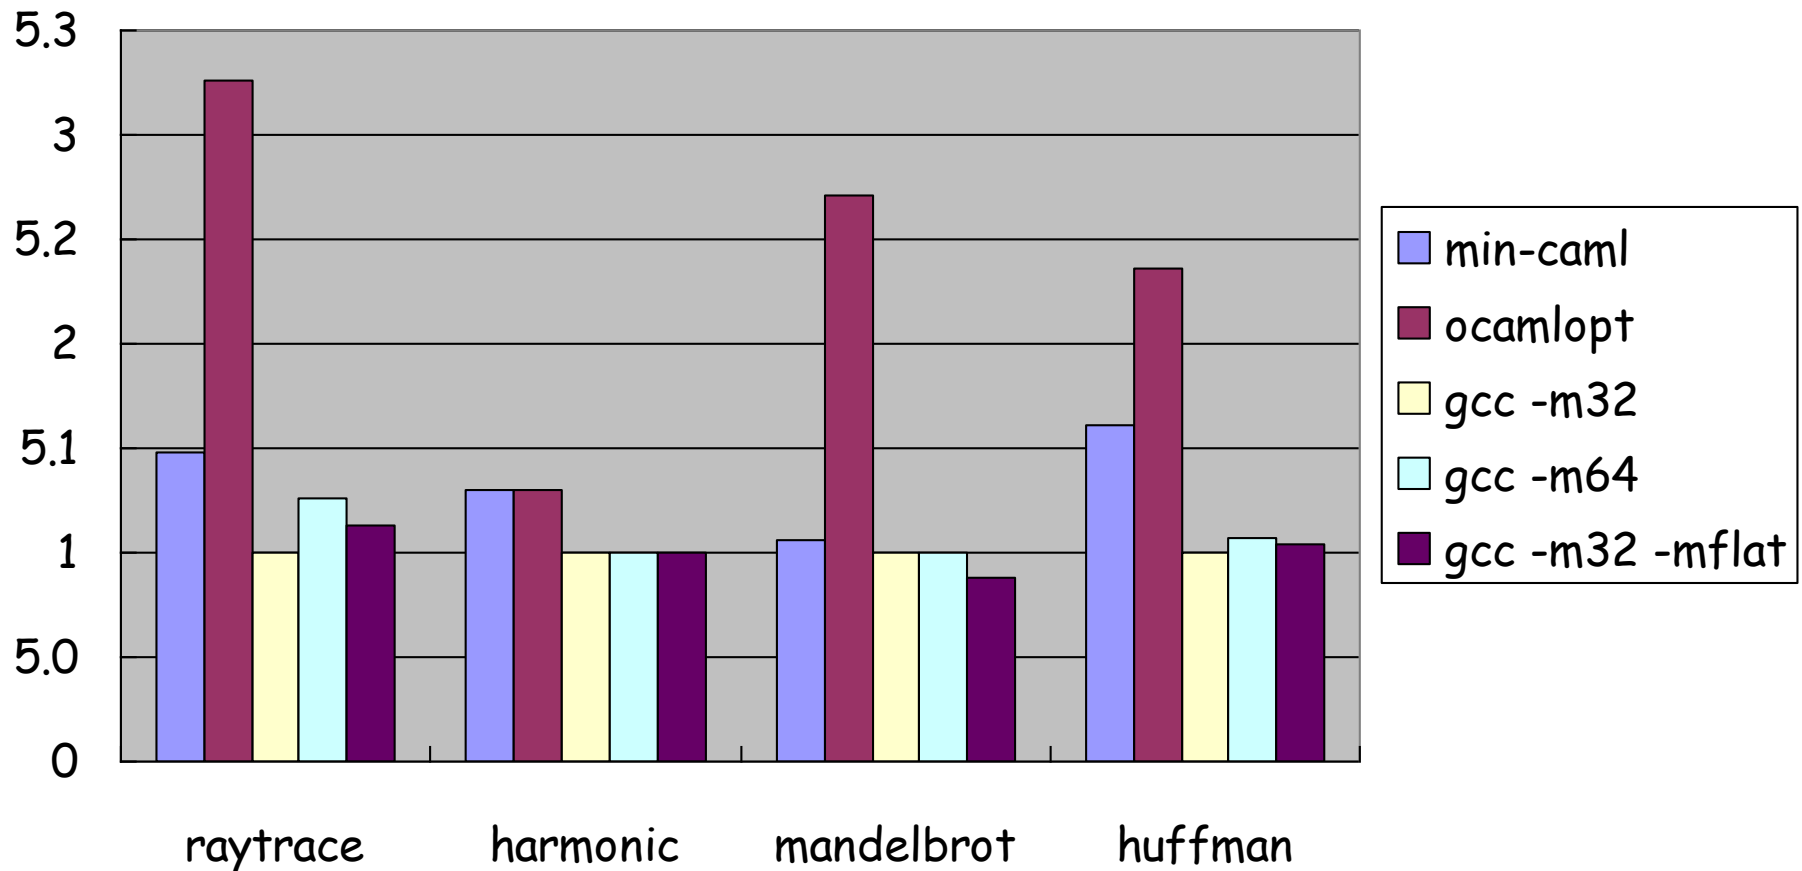
Applications

- Functional
 - Ackermann
 - Fibonacci
 - Takeuchi
- Imperative
 - Ray tracing
 - Harmonic function
 - Mandelbrot set
 - Huffman encoding

Execution Time of Functional Programs (min-caml = 1)



Execution Time of Imperative Programs (gcc -m32 = 1)



Summary

"Simple and efficient compiler
for a minimal functional language"

Future work:

- Improve the register allocation
 - By far more complex than other modules
 - Too slow at compile time
- Retarget to IA-32
 - 2-operand instructions (which are "destructive" by definition) and FPU stack

<http://min-caml.sf.net/>

